

Introduction to Logistic regression

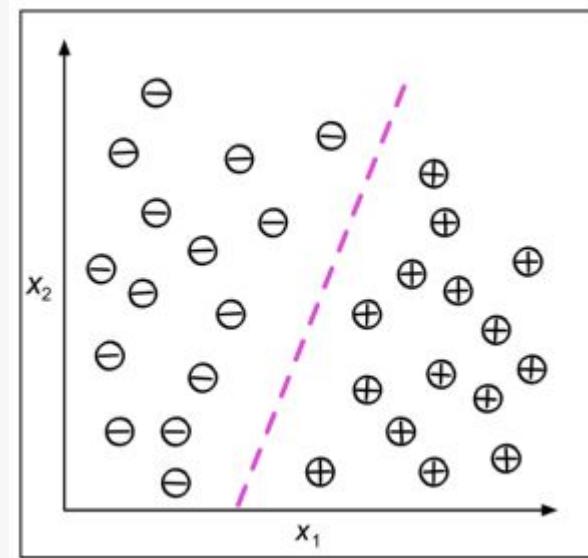
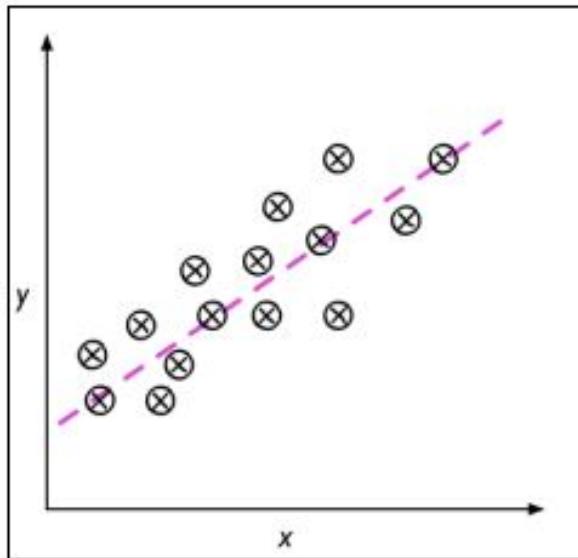


Don't let the name confuse you. For historical reasons it is called 'regression'.

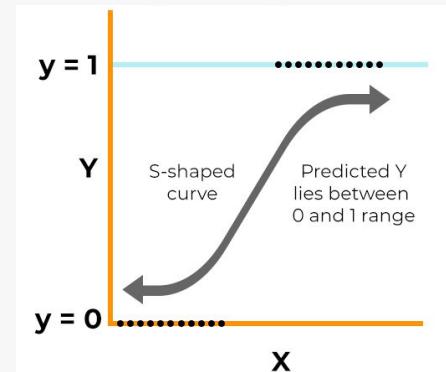
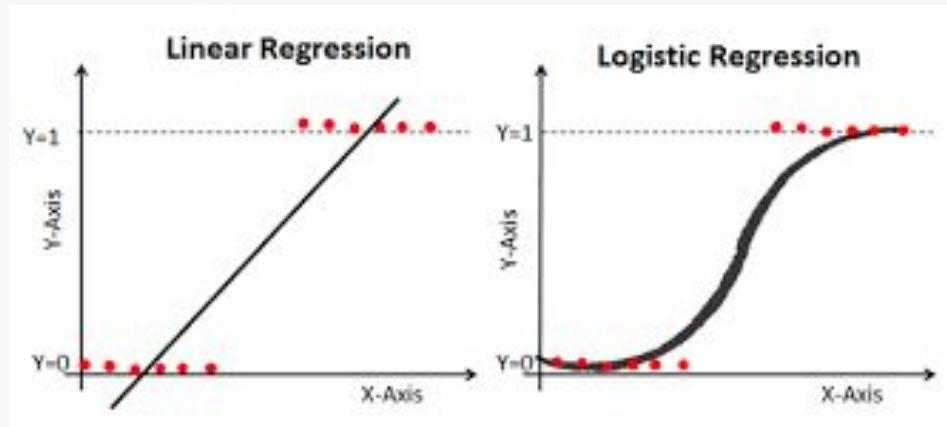
Logistic regression is a CLASSIFICATION algorithm!



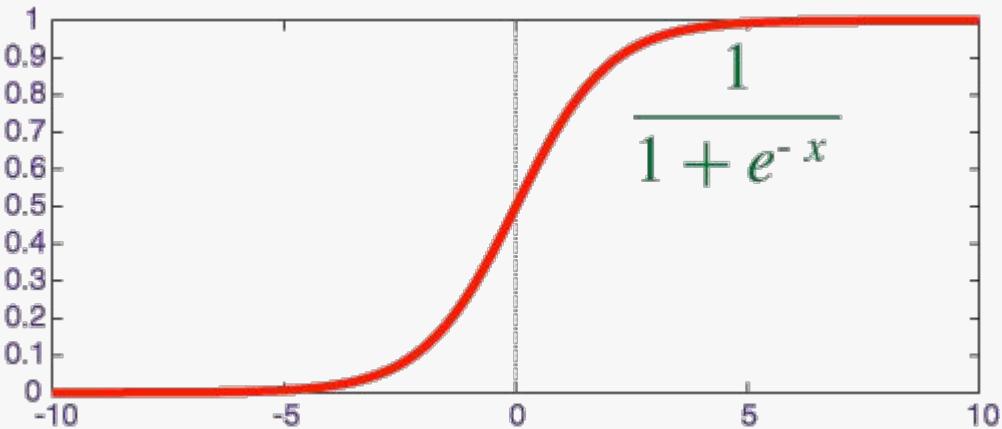
Regression vs Classification



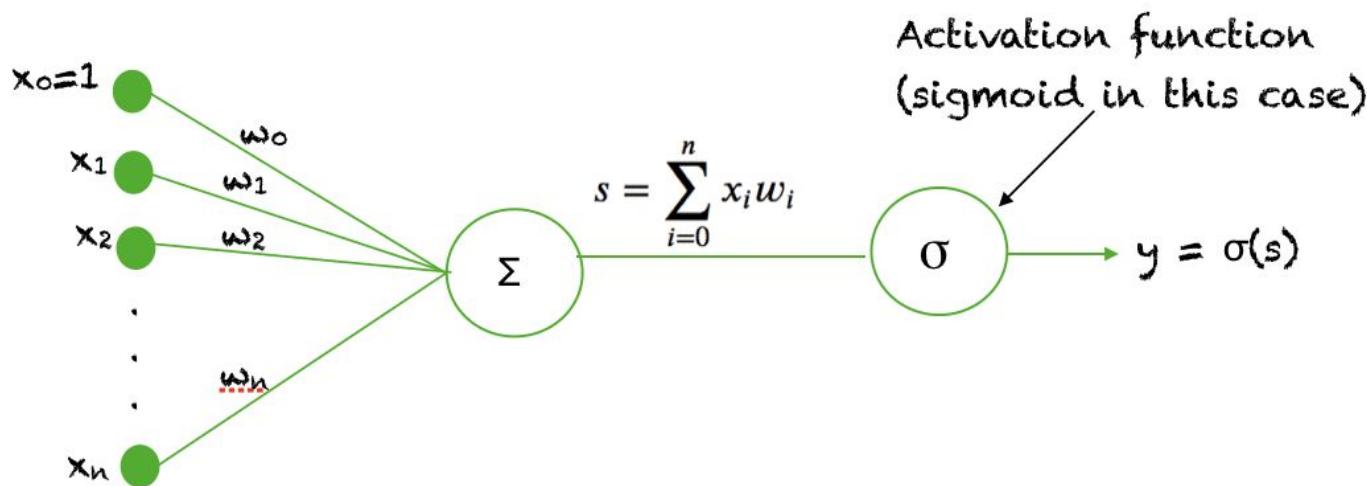
From straight lines to curves



Sigmoid Function



Turning number predictions into class predictions



Cost function for logistic regression

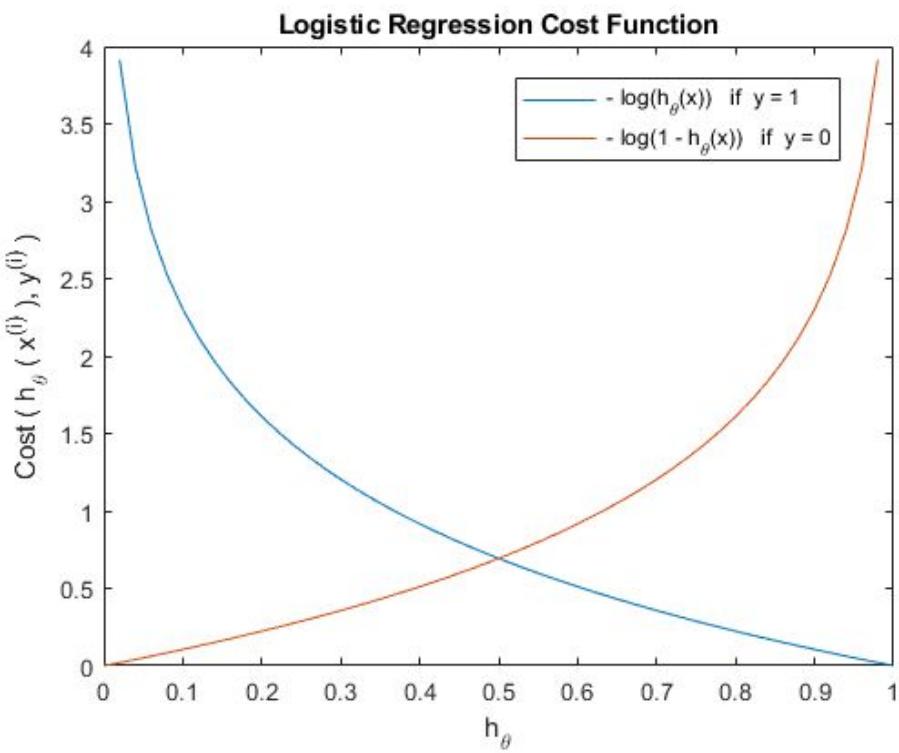


$$J(\theta) = \frac{1}{m} \sum_{i=1}^m Cost(h_\theta(x^{(i)}), y^{(i)})$$

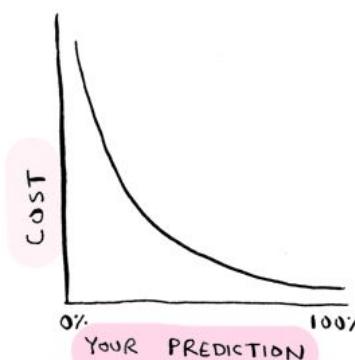
$$J(\theta) = \frac{1}{m} \left[\sum_{i=1}^m -y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)})) \right]$$

m = number of samples

Why logarithm?



IF y is 1...



IF y is 0...



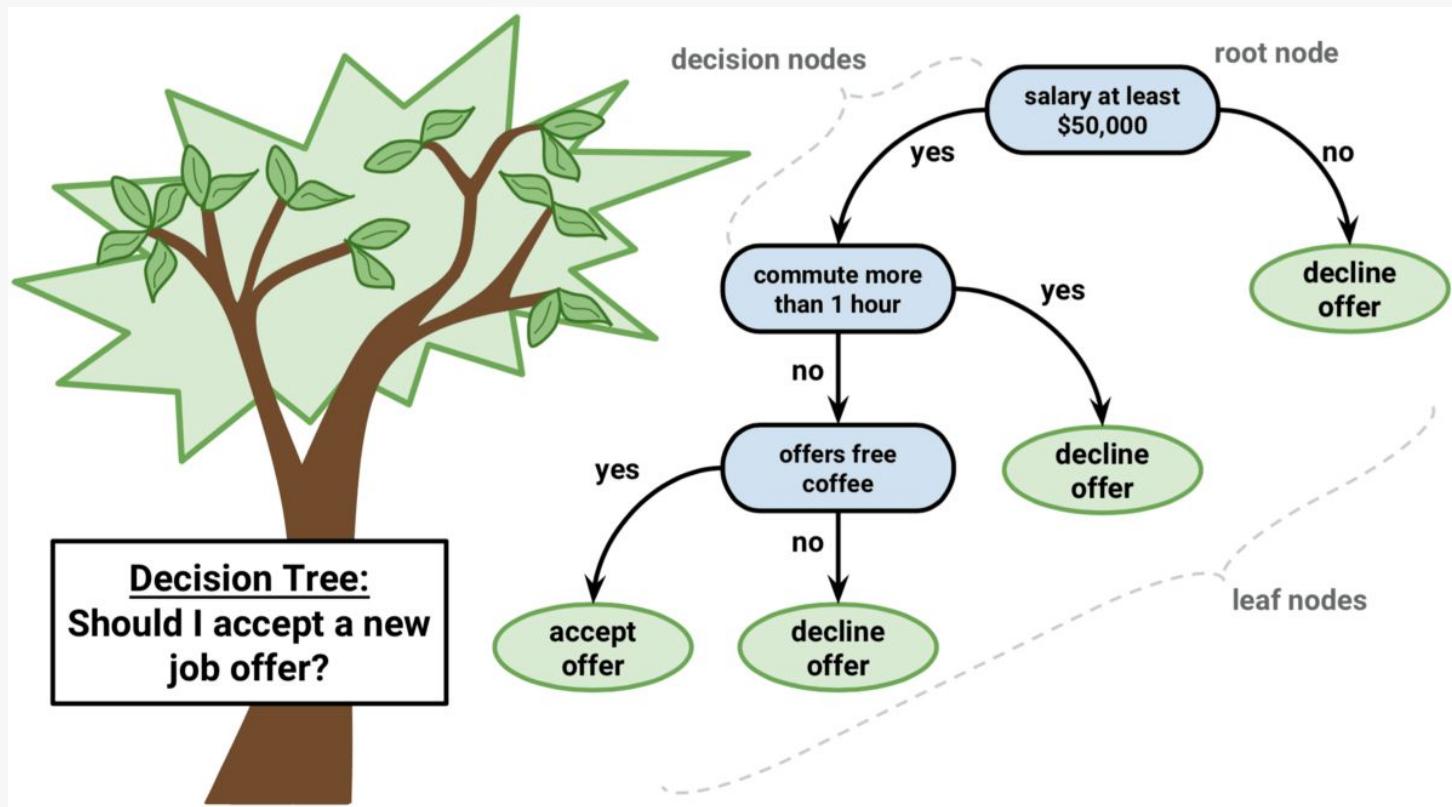
$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y=0 \end{cases}$$

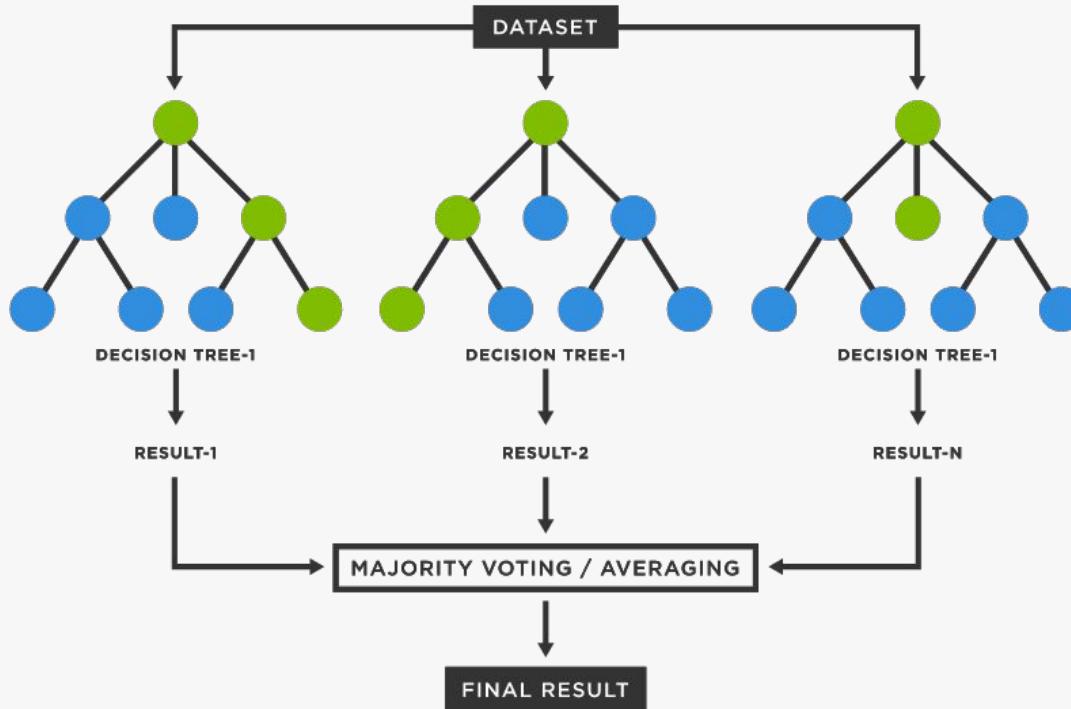
Log loss calculation



Non-Parametric models



Ensemble models



Module 2



Data Science process



Data Engineering Vs Data Science : The number and frequency of experiments

Organizing your experiments

The screenshot shows a Jupyter Notebook interface with the following structure:

- Importing Libraries**: Contains 2 hidden cells.
- Description of the problem**: Contains 1 hidden cell.
- Data Preparation**: Contains 3 hidden cells.
- Model 1 - LR with 1 feature**: Contains 6 hidden cells.
- Model 2 - LR with 2D features**: Contains 10 visible cells with the following code:

```
[ ] In [1]: X1 = np.array([0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, 233, 377, 610])
X2 = np.array([1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, 233, 377, 610, 987])
y = np.array([1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, 233, 377, 610, 987, 1597])
```



Introducing ML communities



- Kaggle
- Huggingface



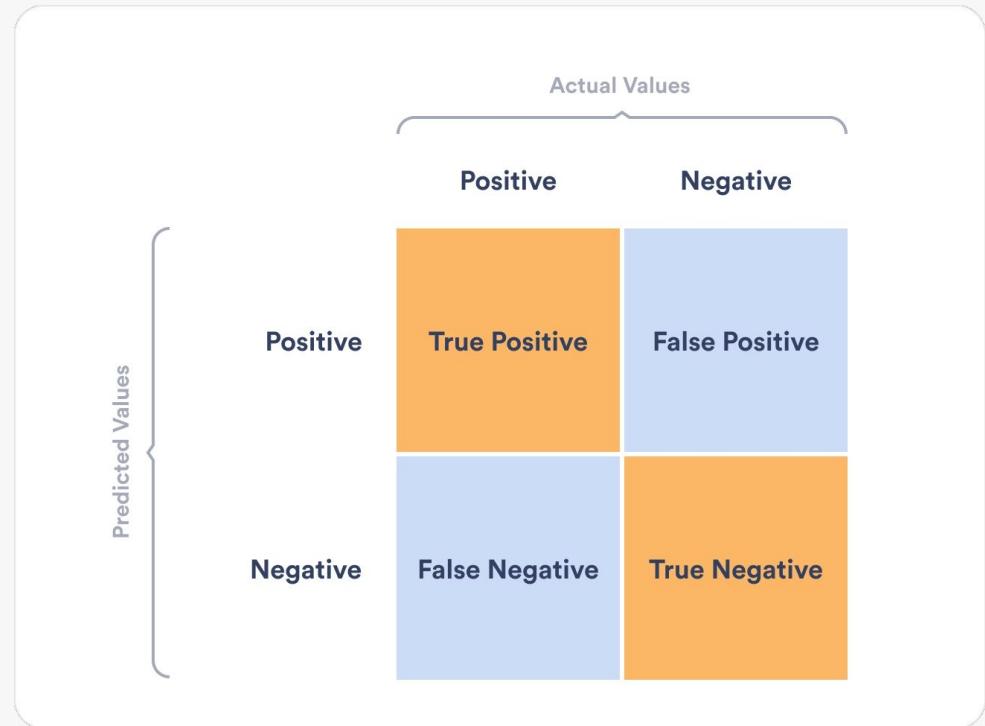
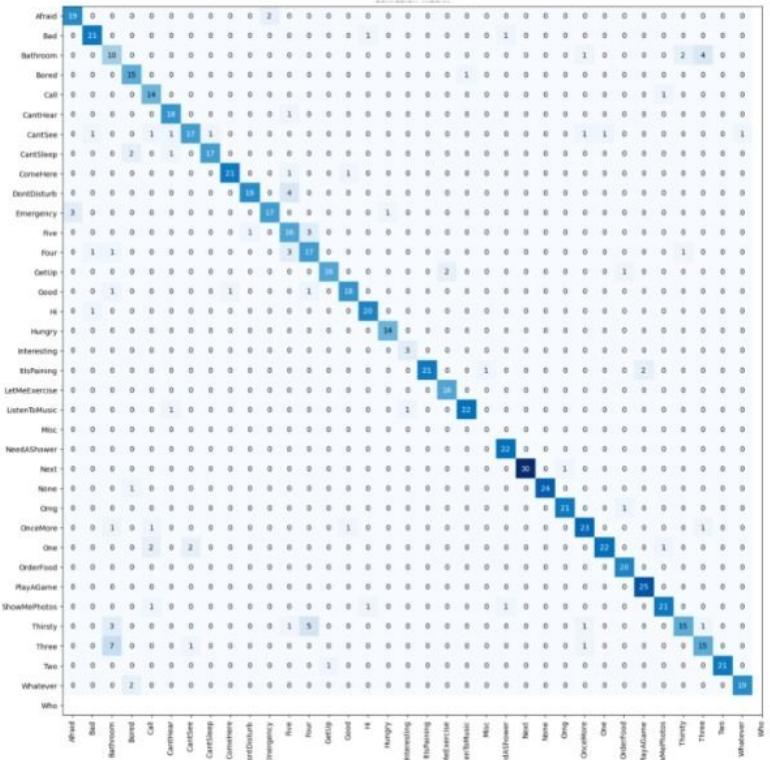
How is my model doing?



Metrics notebook demo



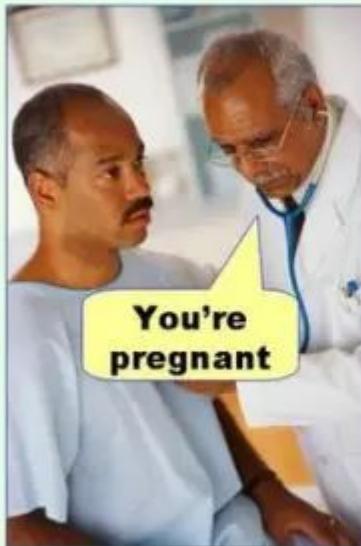
Prediction errors - Confusion Matrix



Metrics - Measuring model performance



Type I error
(false positive)



Type II error
(false negative)



Going beyond accuracy



Accuracy is a limited metric, because models can cheat!!

| Machine learning \ Manual counting | True | False |
|------------------------------------|---------------------|---------------------|
| True | True Positive (TP) | False Positive (FP) |
| False | False Negative (FN) | True Negative (TN) |

Equations:

$$\text{False positive rate (FPR)} = \frac{\text{FP}}{\text{FP} + \text{TN}}$$

$$\text{False negative rate (FNR)} = \frac{\text{FN}}{\text{FN} + \text{TP}}$$

$$\text{Sensitivity} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{Specificity} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

$$\text{Youden index} = \text{Sensitivity} + \text{Specificity} - 1$$

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$



Going beyond *accuracy*

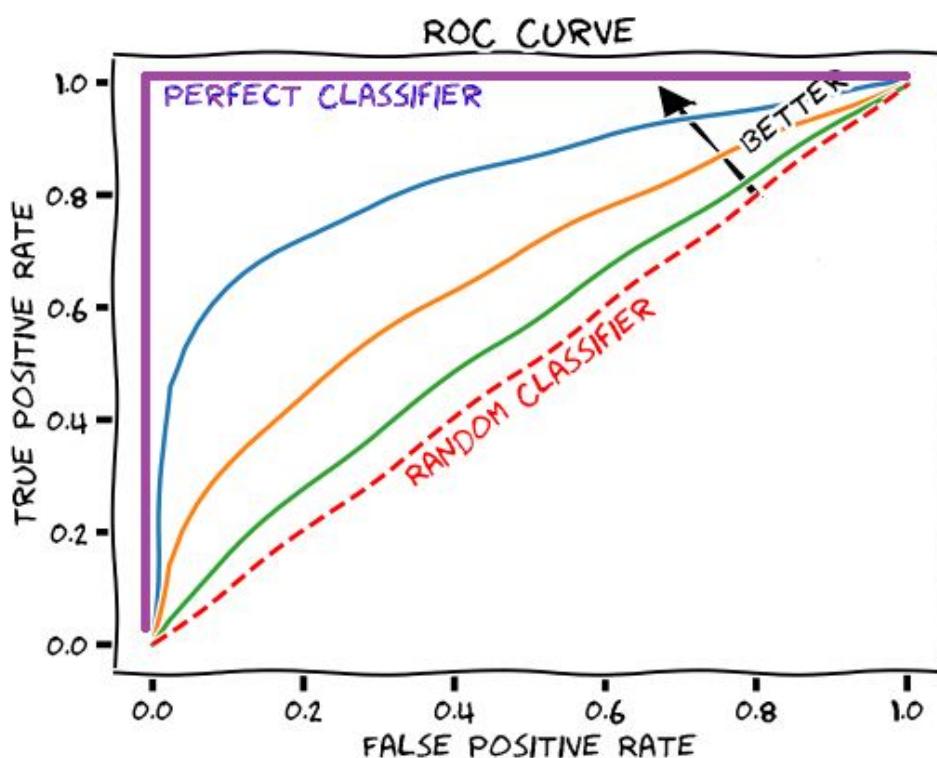


$$\text{Sensitivity} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

$$\text{Specificity} = \frac{\text{True Negatives}}{\text{True Negatives} + \text{False Positives}}$$



AUROC curve



Receiver operating
characteristic
curve.



Introduction to Data Science!



How should we store features?



| ID | First Name | Last Name | Email | Year of Birth |
|----|------------|-----------|--------------------------|---------------|
| 1 | Peter | Lee | plee@university.edu | 1992 |
| 2 | Jonathan | Edwards | jedwards@university.edu | 1994 |
| 3 | Marilyn | Johnson | mjohnson@university.edu | 1993 |
| 6 | Joe | Kim | jkim@university.edu | 1992 |
| 12 | Haley | Martinez | hmartinez@university.edu | 1993 |
| 14 | John | Mfume | jmfume@university.edu | 1991 |
| 15 | David | Letty | dletty@university.edu | 1995 |

Feature Vector

Table: Students

Arrays or vectors!



Vectors and Tensors!



Scalar Vector Matrix Tensor

1

| |
|---|
| 1 |
| 2 |

| | |
|---|---|
| 1 | 2 |
| 3 | 4 |

| | |
|---|---|
| 1 | 2 |
| 1 | 7 |

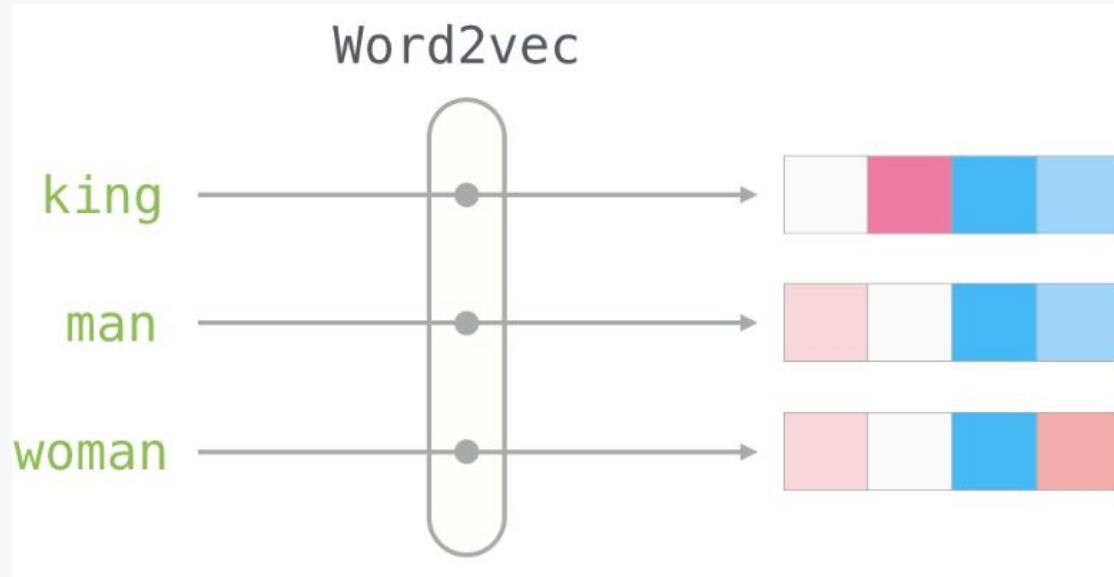
| | |
|---|---|
| 3 | 2 |
| 5 | 4 |



Source: Hadrien Jean's [blog post](#)

Vectors Continued...

Remember the idea of ‘Similarity’ ?

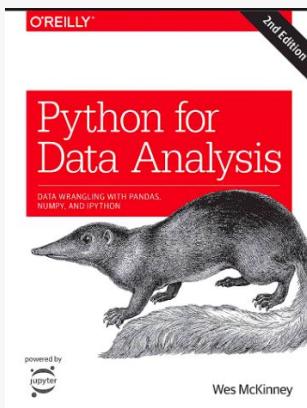


Source: Jay Alammar's [blog post](#)

Open Source Heroes



Pandas: Wes Mckinney



Python: Guido Van Rossum



Jupyter: Fernando Perez



NumPy: Travis Olyphant

