

Module 9



AI Hardware



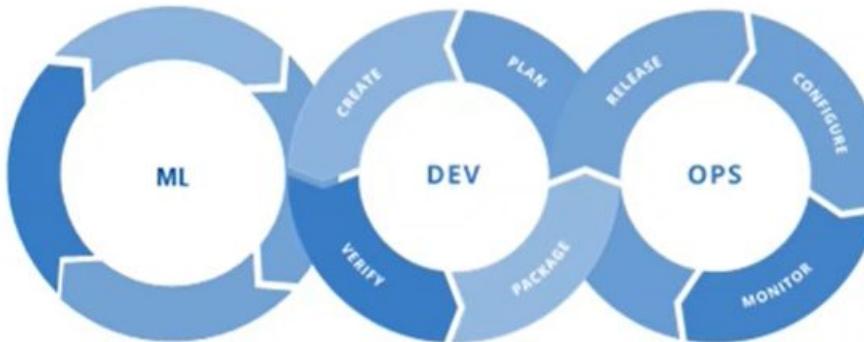
- AWS EC2
- GOOGLE CLOUD
- MICROSOFT AZURE
- JARVIS LABS



From Colab to Deployment



MLOps = ML + DEV + OPS



Experiment

Data Acquisition
Business Understanding
Initial Modeling

Develop

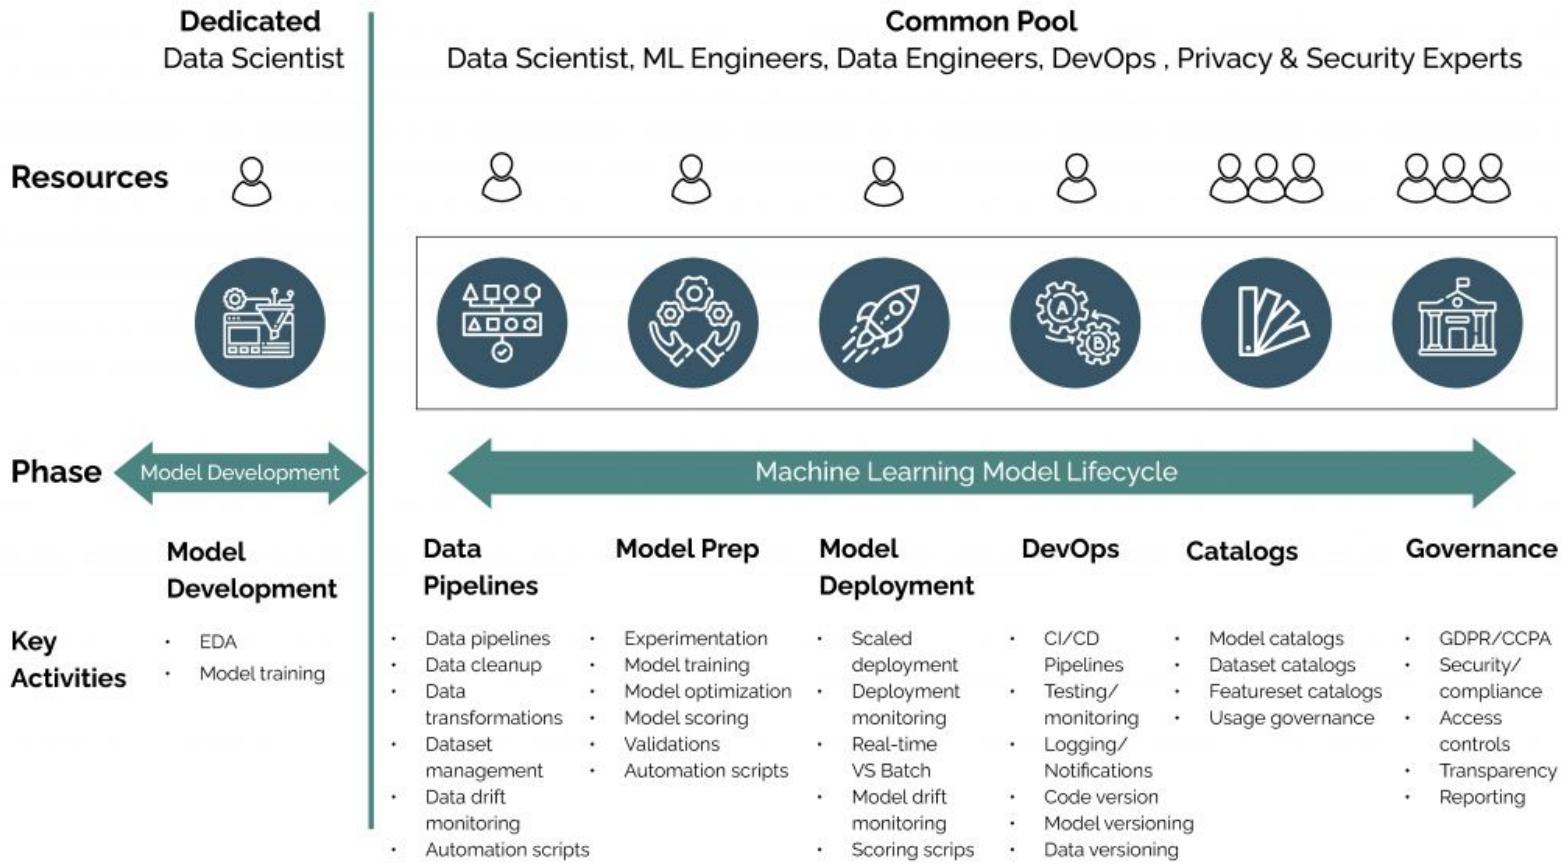
Modeling + Testing
Continuous Integration
Continuous Deployment

Operate

Continuous Delivery
Data Feedback Loop
System + Model Monitoring



Machine Learning Operations (MLOps)



Hardware



Software Tools:

- Libraries like Keras, NumPy, SK-Learn
- Tools like Jupyter notebooks, Google Colab
- Opensource projects like Llama2
- Communities like HF, Kaggle, Github

Hardware:

- GPU???



NVIDIA GPU

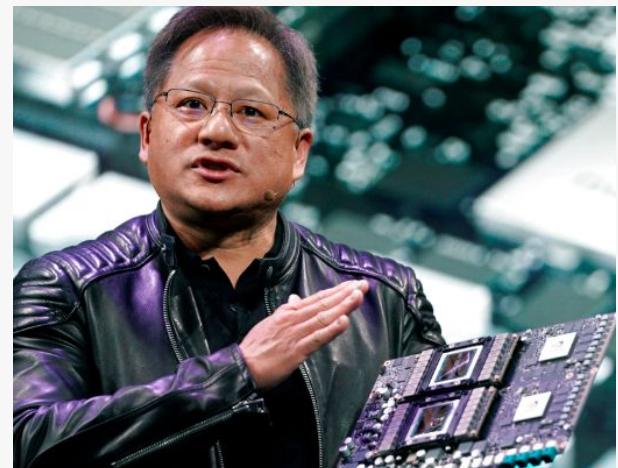


NVIDIA

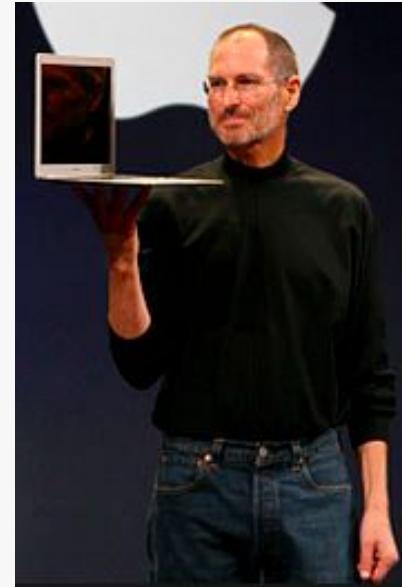
NVIDIA Corp Price % Change



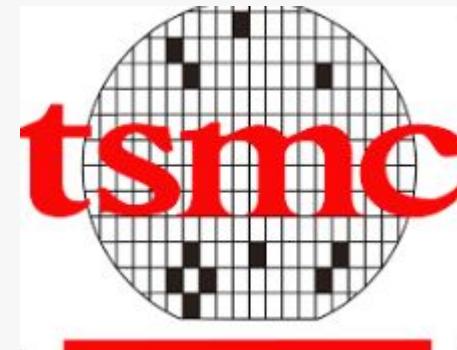
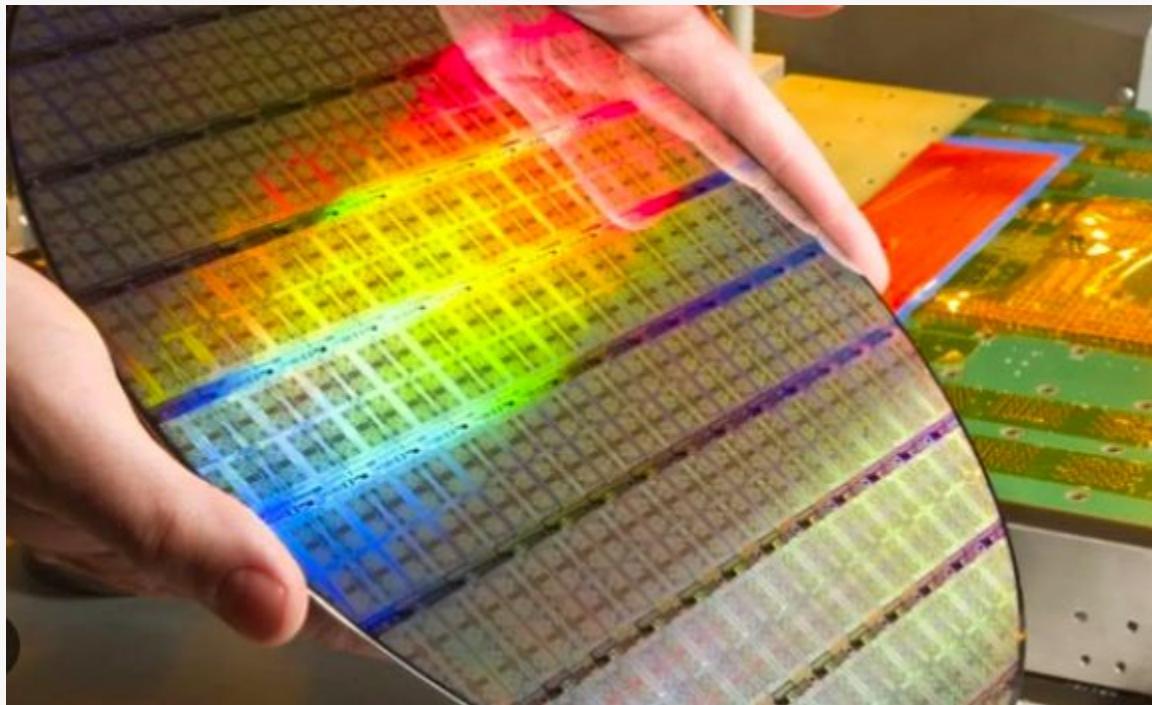
Jensen Huang, CEO of Nvidia Corp



Silicon era heroes



TSMC and Chip Fabrication



EUV - Lithography machine



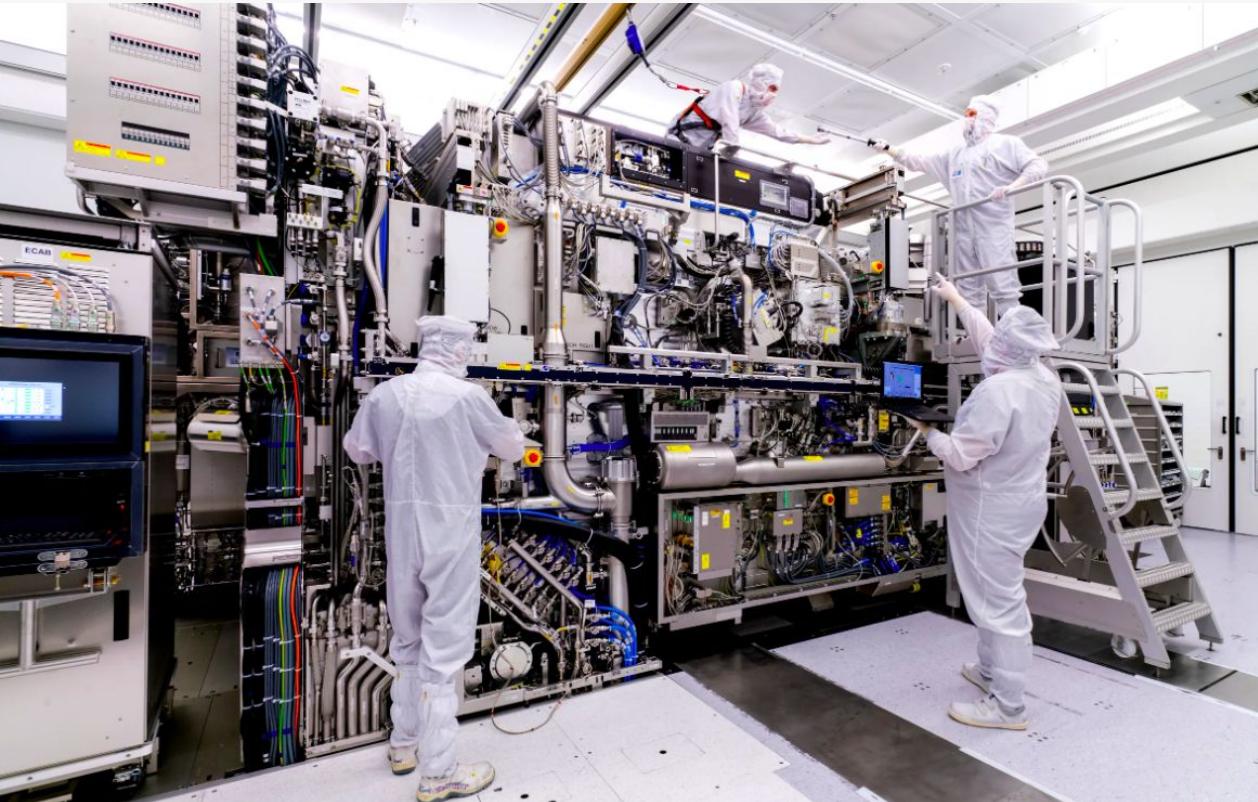
A ball of tin is pulverized by incredibly powerful lasers and explodes into a plasma measuring several times hotter than the surface of the sun.



ASML



ASML - Lithography machine



The GPU revolution of AI

- NVIDIA and CUDA
- TSMC
- ASML
- INTEL & AMD
- ARM PROCESSORS
- APPLE M-Series CHIPS

[Chip War Article: Understanding the GPU revolution of AI](#)

9A. Research Landscape



GOFAI research

Seminal Papers in AI

Arxiv

Research Community - Academia, Industry Labs

Key People and ideas

Conferences

Twitter and debates

How to read ML papers

Bringing your research ideas to life - experimentation and publishing work

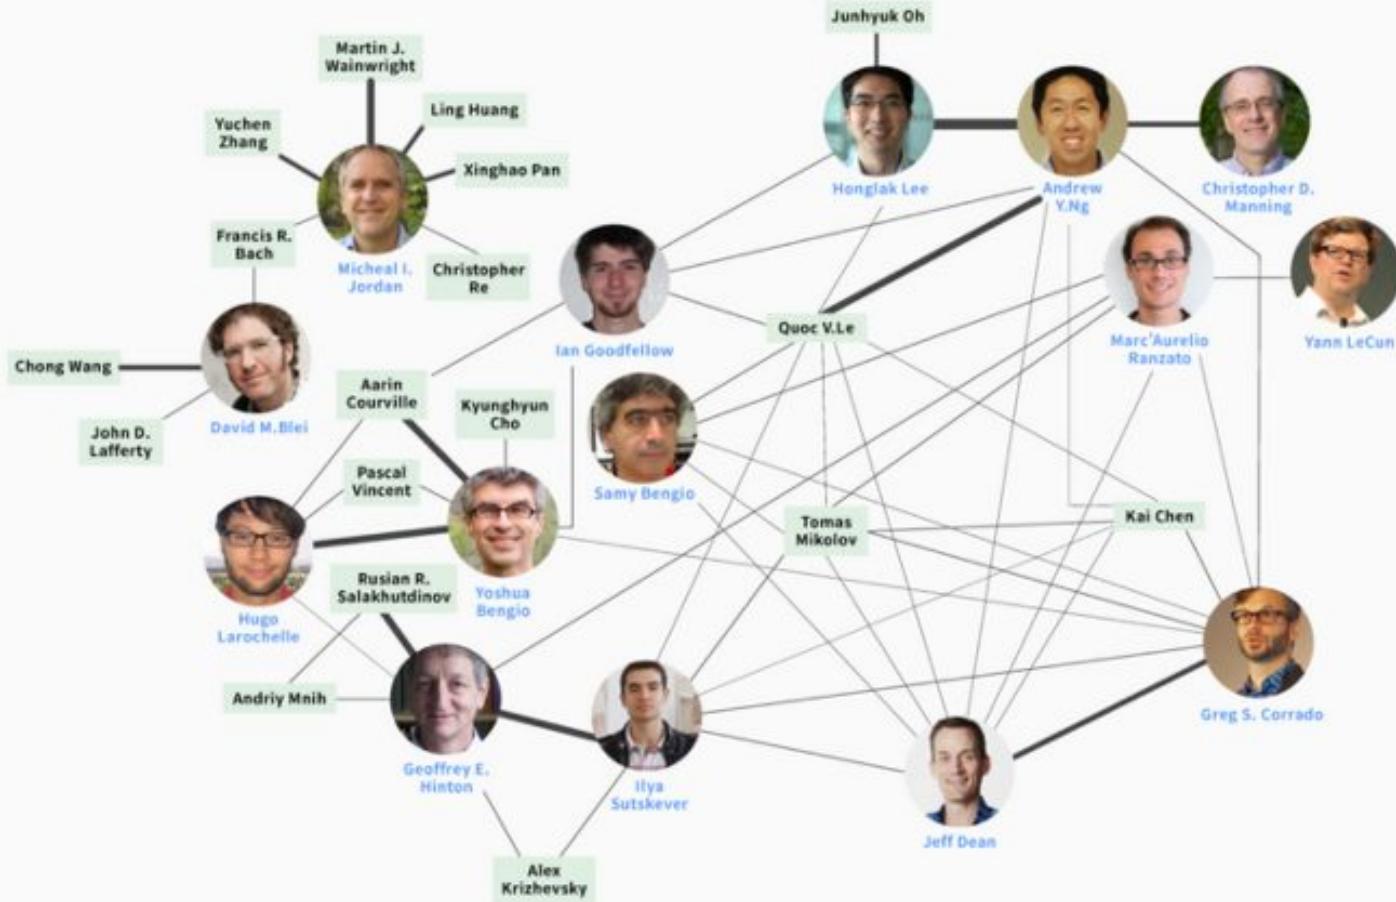
Github / Huggingface





Who knows •the AI Researchers ?





AI Conferences

Neural Information Processing Systems (NeurIPS): One of the largest and most prestigious conferences in machine learning and AI. It covers a wide range of topics including deep learning, reinforcement learning, computer vision, and more.

Conference on Computer Vision and Pattern Recognition (CVPR): A premier conference for computer vision research, covering topics such as image and video analysis, object recognition, and scene understanding.

International Conference on Machine Learning (ICML): A leading conference in machine learning that covers a broad spectrum of topics, including deep learning, optimization, theoretical foundations, and more.

International Joint Conference on Artificial Intelligence (IJCAI): A major AI conference that covers a wide range of AI-related topics, including machine learning, robotics, natural language processing, and more.

Association for Computational Linguistics (ACL): A key conference for natural language processing and computational linguistics, covering areas like language understanding, generation, and translation.

International Conference on Learning Representations (ICLR): A conference focused on the study of deep learning representations and optimization, with an emphasis on novel architectures and techniques.



Labs



9B. Startups



Core:

- Gen AI Startups - Current trends
- Silicon Valley vs Bengaluru
- Creating a product - Levels

Other:

- Building a team
- Fundraising fundamentals
- The startup mindset and the inevitable grind
- GenAI startup opportunities



Module 10



Course Recap



- What have we learnt?
- Student flash presentations
- How can we improve the course?
- How to scale the community?



Capstone Project



Course credits

Ashwin



Gautham



Anuj



Shivansh



Sai



Karthik



Vishnu



Pranav



Smit



Vig

