



RLHF

What, Why and How?

Arvind Nagaraj
Dec 2024

Context - SFT step in LLM training



Prompt *Explain the moon landing to a 6 year old in a few sentences.*

Completion GPT-3

Explain the theory of gravity to a 6 year old.

Explain the theory of relativity to a 6 year old in a few sentences.

Explain the big bang theory to a 6 year old.


Explain evolution to a 6 year old.

InstructGPT

People went to the moon, and they took pictures of what they saw, and sent them back to the earth so we could all see them.



GPT Assistant training pipeline



Step 1

Collect demonstration data, and train a supervised policy.

A prompt is sampled from our prompt dataset.

Explain the moon landing to a 6 year old

A labeler demonstrates the desired output behavior.



Some people went to the moon...

This data is used to fine-tune GPT-3 with supervised learning.



Step 2

Collect comparison data, and train a reward model.

A prompt and several model outputs are sampled.

Explain the moon landing to a 6 year old

A B
Explain gravity... Explain war...
C D
Moon is natural satellite of... People went to the moon...

A labeler ranks the outputs from best to worst.



D > C > A = B

This data is used to train our reward model.




D > C > A = B

Step 3

Optimize a policy against the reward model using reinforcement learning.

A new prompt is sampled from the dataset.

Write a story about frogs



Once upon a time...



r_k

The policy generates an output.

The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.

src



PPO



Proximal

Policy

Optimization



PPO (RL)



RL

HF



REINFORCEMENT LEARNING



Supervised learning = Perception

Reinforcement learning = Judgement



Action and Choices



Robotics and Games




Real time decisions with Mitra

Atari games



Rewards and Penalties



Driving from A to B

← from HSR Layout, Bengaluru, Karnataka to Kempegowda International Airport Bengaluru

1 hr 29 min (53.9 km)

via Outer Ring Rd/Srinagar - Kanyakumari Hwy and NH 44

Best route now, avoids road closure

⚠ This route has tolls.

HSR Layout
Bengaluru, Karnataka

Take 17th Cross Rd, 14th Main Rd and Service Rd to 100 Feet Ring Rd/Outer Ring Rd/Srinagar - Kanyakumari Hwy in Agara Village

4 min (1.5 km)

↑ Head west on 17th Cross Rd toward 17th Main Rd
ⓘ Pass by HDFC Bank (on the right)

700 m

↗ Turn right at 14th Main Rd

130 m

↑ Continue onto 14th Main Rd
ⓘ Pass by Mantri Ceramics (on the right)


240 m

↑ Continue straight onto 14th Main Rd/Jakkasandra Cross Rd

66 m

↗ Turn right onto Service Rd

350 m



Sparse rewards



Rewards are sparse but when you get them, it's usually insightful.



RL is still early



- Explore or exploit?
- Sparse rewards
- Credit assignment?



John Schulman's lecture at Berkeley



Prof.Yoav Goldberg's arguments on : Why RL for LLMs ?



[**yoavg/rl-for-langs.md**](#)



PPO “algorithm”

Algorithm 5 PPO with Clipped Objective

Input: initial policy parameters θ_0 , clipping threshold ϵ

for $k = 0, 1, 2, \dots$ **do**

 Collect set of partial trajectories \mathcal{D}_k on policy $\pi_k = \pi(\theta_k)$

 Estimate advantages $\hat{A}_t^{\pi_k}$ using any advantage estimation algorithm

 Compute policy update

$$\theta_{k+1} = \arg \max_{\theta} \mathcal{L}_{\theta_k}^{CLIP}(\theta)$$

 by taking K steps of minibatch SGD (via Adam), where

$$\mathcal{L}_{\theta_k}^{CLIP}(\theta) = \mathbb{E}_{\tau \sim \pi_k} \left[\sum_{t=0}^T \left[\min(r_t(\theta) \hat{A}_t^{\pi_k}, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t^{\pi_k}) \right] \right]$$

end for



PPO's Clipped surrogate objective function

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right]$$

An extra precaution - KL Divergence



$$\underset{\theta}{\text{maximize}} \quad \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)} \hat{A}_t \right] - \beta \hat{\mathbb{E}}_t [\text{KL}[\pi_{\theta_{\text{old}}}(\cdot | s_t), \pi_{\theta}(\cdot | s_t)]]$$



KL Divergence



$$KL(P||Q) = \sum p_i(x) \log\left(\frac{p_i(x)}{q_i(x)}\right)$$



Finetuning

Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.




One-shot

In addition to the task description, the model sees a single example of the task. No gradient updates are performed.




Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.



Fine-tuning

The model is trained via repeated gradient updates using a large corpus of example tasks.



It is becoming a lot more accessible to finetune LLMs:

- Parameter Efficient FineTuning (PEFT), e.g. LoRA
- Low-precision inference, e.g. bitsandbytes
- Open-sourced high quality base models, e.g. LLaMA


Keep in mind:

- Requires a lot more technical expertise
- Requires contractors and/or synthetic data pipelines
- A lot slower iteration cycle
- SFT is achievable
- RLHF is research territory

[Language Models are Few-Shot Learners, Brown et al. 2020]



RL is Hard!



[Secrets of RLHF in Large Language Models Part I: PPO, Zheng, et al. 2023]

HF TRL library

Step 1: **SFTTrainer**

Train your model on your favorite dataset

```
from trl import SFTTrainer

trainer = SFTTrainer(
    "facebook/opt-350m",
    train_dataset=dataset,
    dataset_text_field="text",
    max_seq_length=512,
)

trainer.train()
```

Step 2: **RewardTrainer**

Train a preference model on a comparison data to rank generations from the supervised fine-tuned (SFT) model

```
from trl import RewardTrainer

trainer = RewardTrainer(
    model=model,
    args=training_args,
    tokenizer=tokenizer,
    train_dataset=dataset,
)

trainer.train()
```

Step 3: **PPOTrainer**

Further optimize the SFT model using the rewards from the reward model and PPO algorithm

```
from trl import PPOConfig, PPOTrainer


trainer = PPOTrainer(
    config,
    model,
    tokenizer=tokenizer,
)

for query in dataloader:
    response = model.generate(query)
    reward = reward_model(response)
    trainer.step(query, response, reward)
```




We still want RL's benefits though!

- Eval: **win rate** against the 175B parameter SFT model (% of responses humans like better)



[Training language models to follow instructions with human feedback, Ouyang, et al., 2022]



DIRECT PREFERENCE OPTIMIZATION



What do we need?



Bradley-Terry Model



$\text{reward_diff} = \log(\text{sigmoid}(\text{reward}_{\text{winner}}) - \text{reward}_{\text{loser}}))$

$\text{LOSS} = \text{maximize}(\text{reward_diff})$

$\text{LOSS} = \text{minimize}(\text{reward_diff} \times -1)$



Bradley-Terry Model



reward_diff = log (sigmoid ((reward_{winner}) - (reward_{loser})))

LOSS = maximize (reward_diff)

LOSS = minimize (reward_diff X -1)

$$\mathcal{L}_R(r_\phi, \mathcal{D}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} [\log \sigma(r_\phi(x, y_w) - r_\phi(x, y_l))]$$



Direct Preference Optimization: Simplifying RLHF

RLHF Objective

(get high reward, stay close to reference model)

$$\max_{\pi} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi(y|x)} [r(x, y)] - \beta \mathbb{D}_{\text{KL}}(\pi(\cdot | x) \| \pi_{\text{ref}}(\cdot | x))$$

any reward function



Closed Form solution



Basically, An ideal mathematical solution (*may not be computable!*)

$$\text{Optimal Policy} = \frac{\text{Reference Policy} \times \text{assigned reward}}{\text{Normalization Value}}$$

- Normalization treats all sequence lengths with equal fairness



Direct Preference Optimization: Simplifying RLHF

RLHF Objective

(get high reward, stay close to reference model)

$$\max_{\pi} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi(y|x)} [r(x, y)] - \beta \mathbb{D}_{\text{KL}}(\pi(\cdot | x) \| \pi_{\text{ref}}(\cdot | x))$$

any reward function

Closed-form Optimal Policy

(write **optimal policy** as function of **reward function**; from prior work)

$$\pi^*(y | x) = \frac{1}{Z(x)} \pi_{\text{ref}}(y | x) \exp \left(\frac{1}{\beta} r(x, y) \right)$$

$$\text{with } Z(x) = \sum_y \pi_{\text{ref}}(y | x) \exp \left(\frac{1}{\beta} r(x, y) \right)$$

Slides by : [Eric Mitchell](#)'s(Stanford University) [lecture at Berkeley](#)



Z - the problem!

SOFTMAX FUNCTION

$$s(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$



Closed-form Optimal Policy

(write **optimal policy** as
function of **reward
function**; from prior work)

$$\pi^*(y \mid x) = \frac{1}{Z(x)} \pi_{\text{ref}}(y \mid x) \exp \left(\frac{1}{\beta} r(x, y) \right)$$

with $Z(x) = \sum_y \pi_{\text{ref}}(y \mid x) \exp \left(\frac{1}{\beta} r(x, y) \right)$



Direct Preference Optimization: Simplifying RLHF

RLHF Objective

(get high reward, stay close to reference model)

$$\max_{\pi} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi(y|x)} [r(x, y)] - \beta \mathbb{D}_{\text{KL}}(\pi(\cdot | x) \| \pi_{\text{ref}}(\cdot | x))$$

any reward function

Closed-form Optimal Policy

(write optimal policy as function of reward function; from prior work)

$$\pi^*(y | x) = \frac{1}{Z(x)} \pi_{\text{ref}}(y | x) \exp\left(\frac{1}{\beta} r(x, y)\right)$$

$$\text{with } Z(x) = \sum_y \pi_{\text{ref}}(y | x) \exp\left(\frac{1}{\beta} r(x, y)\right)$$

Note intractable sum over possible responses; can't immediately use this

Rearrange

(write any reward function as function of optimal policy)

$$r(x, y) = \underbrace{\beta \log \frac{\pi^*(y | x)}{\pi_{\text{ref}}(y | x)} + \beta \log Z(x)}_{\text{Some parameterization of a reward function}}$$

Ratio is **positive** if policy likes response more than reference model, **negative** if policy likes response less than ref. model

Direct Preference Optimization: Putting it together

A loss function on
reward functions

+

A transformation
between reward
functions and policies

=

A loss function
on policies



Direct Preference Optimization: Putting it together

A loss function on reward functions

Derived from the Bradley-Terry model of human preferences

$$\mathcal{L}_R(r, \mathcal{D}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} [\log \sigma(r(x, y_w) - r(x, y_l))]$$

+

A transformation between reward functions and policies

Reward function for which that policy is optimal

Any policy (w/ mild assumptions)!

$$r_{\pi_\theta}(x, y) = \beta \log \frac{\pi_\theta(y \mid x)}{\pi_{\text{ref}}(y \mid x)} + \beta \log Z(x)$$

When substituting, the **log Z term cancels**, because the loss only uses **difference in rewards**

=

A loss function on policies

$$\mathcal{L}_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w \mid x)}{\pi_{\text{ref}}(y_w \mid x)} - \beta \log \frac{\pi_\theta(y_l \mid x)}{\pi_{\text{ref}}(y_l \mid x)} \right) \right]$$

Reward of preferred response

Reward of dispreferred response

Since π_θ is normalized, we've lost a degree of freedom, but not expressiveness (see paper)



Gradient mechanics of DPO

What does the gradient of the loss look like? That is,

$$\nabla_{\theta} \mathcal{L}_{\text{DPO}}(\pi_{\theta}; \pi_{\text{ref}}) = ?$$

$$-\beta \mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \left[\sigma(\hat{r}_{\theta}(x, y_l) - \hat{r}_{\theta}(x, y_w)) \left[\nabla_{\theta} \log \pi(y_w \mid x) - \nabla_{\theta} \log \pi(y_l \mid x) \right] \right]$$

Per-example weight: **Higher weight when the reward model is wrong**

Increase the likelihood of the preferred completions

Decrease the likelihood of the dispreferred completions

$$\hat{r}_{\theta}(x, y) = \beta \log \frac{\pi_{\theta}(y \mid x)}{\pi_{\text{ref}}(y \mid x)}$$



TRL implementation

```
Blame 1246 lines (1088 loc) · 60.3 KB

def dpo_loss(
    pi_logratios = pi_logratios.to(self.accelerator.device)
    ref_logratios = ref_logratios.to(self.accelerator.device)
    logits = pi_logratios - ref_logratios


    # The beta is a temperature parameter for the DPO loss, typically something in the range of 0.1 to 0.5.
    # We ignore the reference model as beta → 0. The label_smoothing parameter encodes our uncertainty about
    # calculates a conservative DPO loss.
    if self.loss_type == "sigmoid":
        losses = (
            -F.logsigmoid(self.beta * logits) * (1 - self.label_smoothing)
            - F.logsigmoid(-self.beta * logits) * self.label_smoothing
        )
    elif self.loss_type == "hinge":
        losses = torch.relu(1 - self.beta * logits)
    elif self.loss_type == "ipo":
        # eqn (17) of the paper where beta is the regularization parameter for the IPO loss, denoted by tau
        losses = (logits - 1 / (2 * self.beta)) ** 2
    elif self.loss_type == "kto_pair":
        # eqn (7) of the HALOs paper
        chosen_KL = (policy_chosen_logps - reference_chosen_logps).mean().clamp(min=0)
        rejected_KL = (policy_rejected_logps - reference_rejected_logps).mean().clamp(min=0)

        chosen_logratios = policy_chosen_logps - reference_chosen_logps
        rejected_logratios = policy_rejected_logps - reference_rejected_logps
        # As described in the KTO report, the KL term for chosen (rejected) is estimated using the rejected
        losses = torch.cat(
            (
                1 - F.sigmoid(self.beta * (chosen_logratios - rejected_KL)),
                1 - F.sigmoid(self.beta * (rejected_logratios - chosen_KL)),
            )
        )
    else:
        raise ValueError(f"Unknown loss type: {self.loss_type}")
)

```



Reward-KL trade-off



[Direct Preference Optimization, Rafailov, Sharma, Mitchell, Ermon, Manning, & Finn., NeurIPS 2023]



Direct Preference Optimization: Simplifying RLHF

High-level punchline

If we parameterize our reward model correctly...

...we can extract the optimal policy for our learned reward model **in closed form, with no additional training**

The trick: use a direct correspondence between optimal policy and reward model!

$$\pi(y|x) \Leftrightarrow r(x, y)$$